

# DBMS IMPORTANT QUESTIONS

## Q1)Extended ER feature

Several concepts extend the basic Entity-Relationship (ER) model, allowing for a more detailed and expressive representation of database structures. Here are some of these extended features:

- Weak Entities:** A **weak entity** is a type of entity that does not have its own primary key and cannot be uniquely identified by its own attributes alone<sup>1</sup>. Instead, a weak entity depends on a **strong entity** (also known as a dominant entity) and its primary key for identification<sup>1</sup> .... A weak entity has a **partial key** (represented by a dashed underline) which, combined with the primary key of the identifying strong entity, provides unique identification<sup>2</sup>. Weak entity sets are represented by a double rectangle, and the relationship connecting a weak entity set to its identifying strong entity set is called an **identifying relationship** and is represented by a double diamond<sup>2</sup> ....

- Types of Attributes:** The ER model supports different types of attributes beyond simple attributes that cannot be divided further<sup>3</sup>.

- Multivalued attributes** can hold multiple values for a single entity. These are represented by a **double ellipse** in an ER diagram<sup>4</sup>

- Derived attributes** are attributes whose values can be calculated or derived from the values of other attributes. These are represented by a **dashed ellipse**<sup>4</sup>.

- Relationship Types and Cardinality:** An ER diagram shows the **relationship among entity sets**<sup>5</sup>. The **mapping cardinality** describes the number of entities to which another entity can be associated via a relationship set<sup>6</sup>. For example, the sources illustrate a **many-to-one** relationship between Student and College, where a college can have many students, but a student cannot study in multiple colleges at the same time<sup>7</sup>. The relationship between two strong entity sets is shown by a single diamond symbol, while the relationship between a strong and a weak entity set is shown by a double diamond symbol<sup>3</sup>.

- Generalization and Specialization:** These are techniques used to organize entities into a hierarchy based on their common and distinct features<sup>6</sup>.

- Generalization** is a **bottom-up approach** where common properties of two or more entities are extracted to create a more general, higher-level entity<sup>6</sup>. The original entities become specializations of this generalized entity. For example, STUDENT and FACULTY can be generalized into PERSON with common attributes like P\_NAME and P\_ADD, while specialized attributes like S\_FEE remain in the STUDENT entity<sup>8</sup>.

◦ **Specialization** is a **top-down approach** where a higher-level entity is divided into two or more lower-level sub-entities based on their distinct characteristics<sup>8</sup>. For example, EMPLOYEE can be specialized into DEVELOPER and TESTER, with common attributes like E\_NAME and E\_SAL in EMPLOYEE and specific attributes like TES\_TYPE in TESTER<sup>8</sup>.

• **Aggregation:** Aggregation is used to represent a relationship between an entity and a relationship<sup>9</sup>. In cases where it's necessary to model a relationship involving other relationships, the relationship and its associated entities can be aggregated into a higher-level entity<sup>9</sup>. For example, a relationship 'eval\_for' between 'proj\_guide' and 'evaluation' might need to be considered in relation to 'student', 'project', and 'instructor' combinations, leading to aggregation<sup>9</sup>.

These extended ER features enhance the modeling capabilities, allowing designers to represent more complex real-world scenarios in a database design.

## Q2 Index, Index type.

Based on the sources you provided, here's information about indexes and index types:

### SQL Indexes:

- **Purpose:** Indexes in SQL are created on existing tables to retrieve rows quickly, especially when tables contain thousands of records. They speed up the retrieval of information by first sorting the data on the indexed column(s) and then assigning a ROWID for each row.
- **Creation:** Indexes are created using the CREATE INDEX statement. The basic syntax is CREATE INDEX index\_name ON table\_name (column\_name1, column\_name2...).
- **Unique Indexes:** A unique index ensures data integrity by not allowing any duplicate values to be inserted into the indexed column. They are created using the CREATE UNIQUE INDEX statement: CREATE UNIQUE INDEX index\_name ON table\_name (column\_name1).
- **Implicit Indexes:** These are automatically created by the database server when an object is created. Specifically, indexes are automatically created for **primary key constraints** and **unique constraints**.
- **Explicit Indexes:** These are created using the CREATE INDEX syntax.
- **Single-Column Indexes:** An index created based on only one table column. The syntax is CREATE INDEX index\_name ON table\_name (column\_name).
- **Composite Indexes:** An index on two or more columns of a table. The syntax is CREATE INDEX index\_name ON table\_name (column1, column2). Up to sixteen columns can be defined in an index in Oracle.

- **Considerations:** While indexes speed up data retrieval, they can slow down DML operations like INSERT, UPDATE, and DELETE because both the indexes and tables need to be updated. Indexes should be used on columns that are frequently used for searching. They are generally not needed for small tables or tables with frequent, large batch updates or inserts, or on columns with a high number of NULL values or those that are frequently manipulated.
- **Displaying Index Information:** The SHOW INDEX FROM table\_name command can be used to list all indexes associated with a table.
- **Dropping Indexes:** Indexes can be removed using the DROP INDEX index\_name command.

## MongoDB Indexes:

- **Purpose:** Indexes in MongoDB support the efficient execution of queries. Without indexes, MongoDB would need to scan every document in a collection to find matches, which is inefficient for large datasets.
- **Creation:** In the mongo shell, indexes are created using the ensureIndex() method. The basic syntax is db.COLLECTION\_NAME.ensureIndex({KEY:1}), where KEY is the field to index and 1 represents ascending order; -1 is used for descending order. MongoDB can also automatically create an index on the \_id field.
- **Single Field Indexes:** These are created on a single field. Example:  
db.friends.ensureIndex({ "name" : 1 }).
- **Compound Indexes:** A single index structure that holds references to multiple fields within a collection's documents. Example: db.products.ensureIndex({ item: 1, category: 1, price: 1 }). The sort order (ascending or descending) specified in a compound index can be important for supporting sort operations.
- **Multikey Indexes:** To index a field that holds an array value, MongoDB creates index items for each element in the array. MongoDB automatically determines if a multikey index is needed. At most one field in a compound index can be an array.
- **Indexes on Embedded Fields:** You can create indexes on fields within subdocuments using dot notation. Example: db.people.ensureIndex({ "address.zipcode": 1 }).
- **Indexes on Subdocuments:** You can index an entire subdocument. For equality matches on subdocuments, the field order matters and the subdocuments must match exactly.
- **Unique Indexes:** These enforce a unique constraint on the indexed field, preventing the insertion of documents with duplicate values for that field. They are created using the unique: true option in ensureIndex(): db.collection.ensureIndex({ a: 1 }, { unique: true }). If a collection has existing duplicate values, you can use the dropDups: true option to index the first occurrence and delete subsequent duplicates (note that this will delete data).
- **Sparse Indexes:** Sparse indexes only include documents that contain the indexed field, omitting references to documents that do not have the field. This can save space. They are created using the sparse: true option: db.collection.ensureIndex({ a: 1 }, { sparse: true }). A sparse index with a unique constraint will allow insertion of

documents with unique values for the indexed field or documents that do not include the field. However, it will not permit duplicate values for the indexed field among the documents that have it.

- **Hashed Indexes:** These compute a hash of the value of a field and index the hashed value, supporting equality queries on the hashed field. They are created by specifying hashed as the index type: `db.collection.ensureIndex({_id: "hashed"})`. Hashed indexes support single fields but not multikey indexes.
- **Index Names:** MongoDB assigns a default name to an index based on the indexed keys and their sort order (e.g., `item_1_quantity_-1`). You can also specify a name using the name option in `ensureIndex()`.
- **Viewing Indexes:** The `getIndexes()` method can be used to view the indexes on a collection.

### Q3 Define sequence, and alternate of it.

A **sequence** is a database object that **automatically generates unique numbers**<sup>1</sup> .... It is typically used to **create a primary key value** which must be unique for each row in a table<sup>1</sup> . Sequences can **speed up the efficiency of accessing sequence values when cached in memory**<sup>1</sup> . They can also be a **time-saving object** because they can reduce the amount of application code needed to write a sequence-generating routine<sup>2</sup> . Sequence numbers are stored and generated independently of tables<sup>2</sup> .

The basic syntax for creating a sequence is:

```
CREATE SEQUENCE sequence_name  
[START WITH start_n  
[INCREMENT BY n]  
[ { MAXVALUE maximum_n | NOMAXVALUE } ]  
[ { MINVALUE minimum_n | NOMINVALUE } ]  
[ { CYCLE | NOCYCLE } ]  
[ { CACHE cache_n | NOCACHE } ]  
[ { ORDER | NOORDER } ];
```

#### Alternate to Sequence

If your DBMS doesn't support sequences, or you prefer a different method, you can use:

#### 1. Auto Increment (MySQL / SQL Server Identity Column):

In databases like MySQL and SQL Server, instead of using a sequence object, you can use an **auto-increment column**:

```
sql
CopyEdit
CREATE TABLE students (
  id INT AUTO_INCREMENT PRIMARY KEY,
  name VARCHAR(100)
);
```

## Q4 Natural Join

A **natural join** is an operation used to combine rows from two or more tables based on columns with the same name and data type in those tables. The join condition is implicitly defined by these common columns.

Here's a breakdown of what the sources say about natural join:

- **Basic Concept:** When a relation  $R$  is decomposed into sub-relations  $R_1, R_2, \dots, R_n$ , a lossless join decomposition ensures that when these sub-relations are joined back using the natural join operator (represented by  $\bowtie$ ), the result is the original relation  $R$ . This means no information is lost during the decomposition. The equation for a lossless join decomposition is given as:  $R_1 \bowtie R_2 \bowtie R_3 \dots \bowtie R_n = R$ .
- **How it Works:** The natural join compares each row of the first table with each row of the second table. If the values of all common attribute names are equal, the join operation produces a new row containing the common attributes and the remaining unique attributes from both tables.
- **Lossy vs. Lossless:** If the join of the decomposed relations results in extra tuples (more than the original relation), it is called a **lossy join decomposition**. In this case,  $R_1 \bowtie R_2 \bowtie R_3 \dots \bowtie R_n \supset R$ . The goal of good database design, especially during normalization, is to achieve lossless-join decompositions.

- **Example of Lossless Join:** Consider a relation  $R(A, B, C)$ :

A	B	C
1	2	1
2	5	3
3	3	3

Decomposed into  $R_1(A, B)$  and  $R_2(B, C)$ :

R1(A, B):

A B

1 2

2 5

3 3

R2(B, C):

B C

2 1

5 3

3 3

The natural join of R1 and R2 ( $R1 \bowtie R2$ ) results in:

A B C

1 2 1

2 5 3

3 3 3

This is the same as the original relation R, so this is a lossless join decomposition.

- **Example of Lossy Join:** Consider the same original relation R(A, B, C) decomposed into R1(A, C) and R2(B, C):

R1(A, C):

A C

1 1

2 3

3 3

R2(B, C):

B C

2 1

5 3

3 3

The natural join of R1 and R2 ( $R1 \bowtie R2$ ) results in:

A B C

1 2 1

2 5 3

2 3 3

3 3 3

This relation contains an extra tuple (2 3 3) which was not in the original relation R, making it a lossy join decomposition.

- **Vertical Fragmentation:** Natural join is also crucial in reconstructing a relation that has been vertically fragmented. If a relation  $r(R)$  is vertically fragmented into  $R_1, R_2, \dots, R_n$  such that  $R = R_1 \cup R_2 \cup \dots \cup R_n$ , then the original relation  $r$  can be reconstructed by taking the natural join of its fragments:  $r = r_1 \bowtie r_2 \bowtie r_3 \bowtie \dots \bowtie r_n$ . To ensure this lossless join, a common candidate key (or superkey) is usually included in each of the  $R_i$ , or a special tuple-id attribute might be added to each schema to serve as a candidate key.

The **natural join** is a fundamental relational algebra operation that combines tables based on common attributes. It is a key concept in understanding how to decompose relations effectively and how to reconstruct them without losing information, which is essential for **lossless-join decomposition** in database normalization.

## Q5 How implement full join.

A **FULL JOIN** returns rows when there is a match in **one of the tables** being joined. It essentially combines the results of both a **LEFT JOIN** and a **RIGHT JOIN**.

Here's how it works and how to implement it:

- **Syntax:** The general SQL syntax for a full join is:
  - SELECT column\_name(s)
  - FROM table\_name1
  - FULL JOIN table\_name2
  - ON table\_name1.column\_name = table\_name2.column\_name;

Here, table\_name1 and table\_name2 are the tables you want to join, and column\_name is the column that is used for the join condition. The ON clause specifies how the rows from the two tables should be matched.

- **Behavior:**
  - It returns all rows from the **left table** (table\_name1). If there is a match in the right table (table\_name2), the corresponding right table columns are displayed. If there is no match, the right table columns will contain NULL values.
  - It also returns all rows from the **right table** (table\_name2). If there is a match in the left table (table\_name1), the corresponding left table columns are displayed. If there is no match, the left table columns will contain NULL values.

- Therefore, a full join includes all rows from both tables, filling in NULLs for columns where there isn't a corresponding match in the other table.
- Example:** Consider two tables, "Persons" and "Orders":

**The "Persons" table:**

P_Id	LastName	FirstName	Address	City
1	ABC	A	ST1	pune
2	PQR	B	ST2	pune
3	XYZ	C	ST3	mumbai

**The "Orders" table:**

O_Id	OrderNo	P_Id
1	77895	3
2	44678	3
3	22456	1
4	24562	1
5	34764	15

To list all persons and their orders, and all orders with their persons, you would use the following FULL JOIN statement:

```
SELECT Persons.LastName, Persons.FirstName, Orders.OrderNo
FROM Persons
FULL JOIN Orders
ON Persons.P_Id = Orders.P_Id
ORDER BY Persons.LastName;
```

The result-set would look like this:

LastName	FirstName	OrderNo
ABC	A	22456
ABC	A	24562
PQR	B	
XYZ	C	77895
XYZ	C	44678
		34764

As you can see, all rows from "Persons" and all rows from "Orders" are included. Where a person has no matching order (like 'PQR') or an order has no matching person (like order number '34764' with P\_Id 15), NULL values are present in the respective columns.

- Outer Join Classification:** The sources also indicate that **FULL JOIN** is a type of **SQL Outer Join**. Other types of outer joins include **LEFT JOIN** and **RIGHT JOIN**.
- RDBMS Variations:** It's worth noting that the syntax for outer joins, including full outer joins, can have slight variations depending on the specific Relational



Database Management System (RDBMS) you are using. Some systems might use the explicit FULL OUTER JOIN keyword, while others might have different notations. However, the core concept and functionality remain the same.

## Q6 pl/sql architecture(components,block diagram,engine,database server)

The PL/SQL architecture consists of **three main components**: the **PL/SQL block**, the **PL/SQL Engine**, and the **Database Server**.

### 1. PL/SQL Block:

- This is the section that contains the **actual PL/SQL code**.
- All PL/SQL units are treated as PL/SQL blocks, which serve as the primary input to the architecture.
- A PL/SQL block is logically divided into different sections:
  - **Declarations Section**: This **optional** section starts with the keyword DECLARE and is used to define all variables, cursors, subprograms, and other elements that will be used in the program. If no declarations are needed, this section can be skipped. This must be the first section in a PL/SQL block if present.
  - **Executable Commands Section**: This is a **mandatory** section enclosed between the keywords BEGIN and END. It contains the executable PL/SQL statements of the program and must have at least one executable line of code (even a NULL command). This section can contain both PL/SQL and SQL code.
  - **Exception Handling Section**: This **optional** section starts with the keyword EXCEPTION and contains handlers for runtime errors (exceptions) that may occur in the program. It is the last part of the PL/SQL block, and control cannot return to the execution block from here.

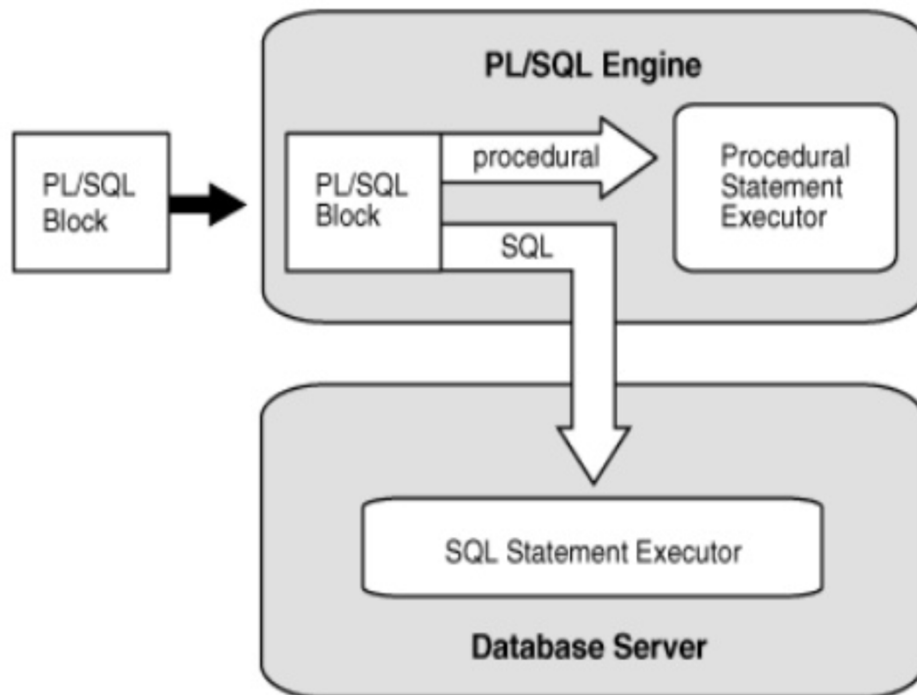
### 2. PL/SQL Engine:

- This component is responsible for the **actual processing of the PL/SQL code**.
- The PL/SQL Engine **separates the PL/SQL units from the SQL parts** within the input.
- The separated PL/SQL units are handled by the PL/SQL Engine itself.
- The **SQL parts are sent to the Database Server** for interaction with the database.
- The PL/SQL Engine can be installed in both the **Database Server** and in an **Application Server**.

### 3. Database Server:

- This is the crucial component that **stores the data**.
- The **PL/SQL Engine** uses the **SQL** from the PL/SQL units to interact with the **Database Server**.
- The Database Server includes an **SQL executor** that parses and executes the input SQL statements.

## PL/SQL Architecture



### Block Diagram (Conceptual Description):

1. A PL/SQL block containing code (declarations, executable commands, and optionally exception handling) is submitted as input.
2. The **PL/SQL Engine** receives this block.
3. The Engine **separates the PL/SQL procedural statements from the SQL statements** within the block.
4. The **PL/SQL procedural statements are processed directly by the PL/SQL Engine**.
5. The **SQL statements are sent by the PL/SQL Engine to the Database Server**.

6. The **SQL executor within the Database Server** parses and executes the **SQL statements**, interacting with the stored data.
7. The results of the SQL execution are sent back to the PL/SQL Engine.
8. The PL/SQL Engine continues processing the PL/SQL block, potentially using the results from the Database Server.
9. Finally, the outcome of the PL/SQL block execution is returned.

In essence, the PL/SQL block provides the instructions, the PL/SQL Engine interprets and executes the procedural parts and delegates SQL operations, and the Database Server manages the data and processes the SQL requests. This separation of concerns allows for efficient execution of database-centric applications.

## Q7 sql vs plsql

Feature	SQL	PL/SQL
Purpose/Usage	Aimed to <b>store, manipulate data</b> stored in relational databases. Used for <b>inputting, modifying, and dropping data</b> . Used for specifying the <b>database schema</b> . Primarily used for <b>data manipulation (DML)</b> and <b>data definition (DDL)</b> operations. Used to <b>retrieve data</b> .	Procedural extension for SQL in Oracle. Combines SQL with <b>procedural features of programming languages</b> . Used to write <b>program blocks/procedures/functions</b> . Used to create <b>applications</b> . Allows instructing the compiler 'what to do' through SQL and 'how to do' through its procedural way.
Unit of Execution	Executes as a <b>single query</b> . DML statements are broken down into instructions for the storage manager.	Executes as a <b>whole block of code</b> . PL/SQL programs are divided and written in <b>logical blocks of code</b> .
Declarative vs. Procedural	<b>Declarative</b> , defines <b>what</b> needs to be done, rather than how. You specify the data you want, and the database figures out how to get it.	<b>Procedural</b> , defines <b>how</b> things need to be done. Provides more control through the use of <b>loops, conditions, and object-oriented concepts</b> . PL/SQL <b>engine separates PL/SQL units and SQL parts</b> . PL/SQL units are handled by the engine, and the <b>SQL parts are sent to the Database Server</b> . This can reduce network traffic by sending an entire block at once.
Interaction with Database Server	Interacts directly with the <b>Database Server</b> to perform DML and DDL operations. SQL statements are sent to the server for execution.	<b>Is an extension of SQL</b> , so it <b>can contain SQL inside it</b> . SQL statements are used to interact with the database server from within PL/SQL blocks.
Contains the Other	<b>Cannot contain PL/SQL code</b> in it.	Implementing <b>stored procedures and functions</b> , writing <b>triggers</b> , handling <b>exceptions (runtime errors)</b> , creating database applications, automating tasks, enforcing complex business rules.
Main Use Cases	Retrieving specific data, inserting new data, updating existing data, deleting data, creating and modifying database objects like tables, views, and indexes.	

## Q8 diff between procedure and function

Both procedures and functions are **named PL/SQL blocks** and are considered **subprograms**. They can be created at the schema level as standalone subprograms using CREATE PROCEDURE or CREATE FUNCTION statements. They can also be part of packages or defined within a PL/SQL block. Both can accept parameters using IN, OUT, or IN OUT modes.

The **significant difference** between a procedure and a function is that **a function must always return a single value**, whereas **a procedure may or may not return a value directly**.

Here's a more detailed comparison:

- **Return Value:**
  - **Function: Must return a single value.** This is specified in the RETURN clause in the function header and the executable section. The datatype of the return value must be declared.
  - **Procedure: Does not necessarily return a value directly.** It primarily performs an action. Procedures can return values to the calling program using OUT or IN OUT parameters.
- **Main Purpose:**
  - **Function: Mainly used to compute and return a value.** It is designed to perform a calculation and give back a result.
  - **Procedure: Mainly used to perform an action.** It can execute a series of PL/SQL and SQL statements to modify data or perform specific tasks.
- **Syntax:**
  - **Procedure:** Created using CREATE [OR REPLACE] PROCEDURE procedure\_name [(parameter\_name [IN | OUT | IN OUT] type [, ...])] [IS | AS] BEGIN <procedure\_body> END [procedure\_name];.
  - **Function:** Created using CREATE [OR REPLACE] FUNCTION function\_name [(parameter\_name [IN | OUT | IN OUT] type [, ...])] RETURN return\_datatype [IS | AS] BEGIN <function\_body> END [function\_name];. Notice the mandatory RETURN return\_datatype clause in the function syntax.
- **Usage:**
  - **Function:** Because it returns a value, a function's output **needs to be assigned to a variable or can be used directly in a SELECT statement**.
  - **Procedure:** Procedures are typically called using the EXECUTE keyword or by simply calling their name from within a PL/SQL block.

In essence, think of **functions as similar to mathematical functions** that take input and produce a specific output, while **procedures are more like a sequence of steps** that perform a task.

## Q9 Types of index and cursor

### Types of Indexes:

According to the sources, there are several types of SQL indexes:

- **Implicit Indexes:** These are indexes that are **automatically created by the database server when an object is created**. Specifically, indexes are automatically created for **primary key constraints and unique constraints**.
- **Explicit Indexes:** These are indexes that are **created using the CREATE INDEX syntax**.

Within explicit indexes, several subtypes are mentioned:

- **Single-Column Indexes:** An index created based on **only one table column**. The syntax is `CREATE INDEX index_name ON table_name (column_name)`.
- **Unique Indexes:** These indexes are used for **performance and data integrity**. They **do not allow any duplicate values** to be inserted into the table. The syntax is `CREATE UNIQUE INDEX index_name ON table_name (column_name)`.
- **Composite Indexes:** An index on **two or more columns of a table**. The syntax is `CREATE INDEX index_name ON table_name (column1, column2)`. In Oracle, you can define up to sixteen columns in an index.

MongoDB, discussed in relation to NoSQL databases, also has different index types:

- **Single Field Indexes:** These are created on a single field within a document. The command to create one is `db.COLLECTION_NAME.ensureIndex({KEY:1})`, where KEY is the field name and 1 denotes ascending order (-1 for descending).
- **Compound Indexes:** A single index structure that holds references to **multiple fields within a collection's documents**. You can build one using `db.collection.ensureIndex({ a: 1, b: 1, c: 1 })`. The sort order specified (ascending or descending) matters for compound indexes in terms of supporting sort operations.
- **Multikey Indexes:** To index a field that holds an **array value**, MongoDB creates index items for each item in the array. You do not need to explicitly specify the multikey type; MongoDB automatically determines it. However, in a multikey compound index, **at most one field can hold an array**.
- **Sparse Indexes:** These indexes **omit references to documents that do not include the indexed field**. They can provide space savings for fields only present in some documents. You can create one with a unique constraint using `db.scores.ensureIndex({ score: 1 }, { sparse: true, unique: true })`.

### Types of Cursors:

According to the sources focusing on PL/SQL, there are **two main types of cursors**:

- **Implicit Cursors:** These are **automatically declared by Oracle** every time an SQL statement is executed, as long as an explicit cursor does not exist for that SQL statement. The user is generally **not aware** of this happening and cannot control or process the information in an implicit cursor. An implicit cursor is automatically associated with every **DML (Data Manipulation) statement (UPDATE, DELETE, INSERT)**. The most recently opened implicit cursor is referred to as the **“SQL%” Cursor**, which has attributes like %FOUND, %ISOPEN, %NOTFOUND, and %ROWCOUNT. For implicit cursors, %ISOPEN always returns FALSE because Oracle closes the SQL cursor automatically. Implicit cursors are less efficient than explicit cursors.
- **Explicit Cursors:** These are **defined by the programmer** for any query that returns **more than one row of data**. An explicit cursor must be **declared in the declaration section of the PL/SQL block**. Working with an explicit cursor involves the following steps:
  - **Declaring** the cursor for initializing memory.
  - **Opening** the cursor for allocating memory.
  - **Fetching** the cursor for retrieving the data, one row at a time.
  - **Closing** the cursor to release the allocated memory. Explicit cursors provide **more programmatic control** to the programmer compared to implicit cursors. They are also generally **more efficient**.

The sources also mention the concept of a **cursor FOR loop**, which is a **shortcut to process cursors**. It implicitly opens, fetches, and closes the cursor, and the record is implicitly declared.

## Q10 Explain closure of set of functional dependency

The **closure of a set of functional dependencies (F)** is the **set of all functional dependencies that are logically implied by F**. This means that if a functional dependency holds whenever all the dependencies in F hold, then that functional dependency is part of the closure of F. The closure of F is denoted by **F<sup>+</sup>**. Importantly, **F<sup>+</sup> is always a superset of F**, meaning it includes all the original functional dependencies in addition to those that can be inferred from them.

For example, if your initial set of functional dependencies F contains  $A \rightarrow B$  and  $B \rightarrow C$ , then the functional dependency  $A \rightarrow C$  is logically implied by F through transitivity. Therefore,  $A \rightarrow C$  would be included in the closure of F, denoted as F<sup>+</sup>.

The sources mention that **Armstrong's axioms are the basic inference rules** used to determine functional dependencies on a relational database and can be used to derive additional functional dependencies from the initial set. These rules allow us to conclude which new functional dependencies must also hold true given an initial set of FDs.

Understanding the closure of a set of functional dependencies is important because when a database is updated, the system needs to check if any functional dependencies are violated. Instead of checking all possible functional dependencies, we can work with the original set and the rules of inference to ensure data consistency with respect to all implied dependencies. The concept of a **canonical cover** is also related, as it's a simplified set of functional dependencies that has the same closure as the original set F.

## Q11 Reflexivity rules

According to the **Reflexive Rule (IR1)**, if **Y is a subset of X**, then **X determines Y**. This can be written as:

**If  $X \supseteq Y$  then  $X \rightarrow Y$ .**

Here, X and Y represent sets of attributes in a relation. The rule states that if a set of attributes Y is entirely contained within another set of attributes X, then X functionally determines Y.

**Example provided in the source:**

If  $X = \{a, b, c, d, e\}$  and  $Y = \{a, b, c\}$ , then according to the reflexive rule,  **$X \rightarrow Y$**  (i.e.,  $\{a, b, c, d, e\} \rightarrow \{a, b, c\}$ ) holds true. This is because  $\{a, b, c\}$  is a subset of  $\{a, b, c, d, e\}$ .



Another example mentioned in the context of trivial functional dependencies:

For a table with attributes Employee\_Id and Employee\_Name, **{Employee\_Id, Employee\_Name} → Employee\_Id** is a trivial functional dependency because Employee\_Id is a subset of {Employee\_Id, Employee\_Name}. Similarly, **Employee\_Id → Employee\_Id** and **Employee\_Name → Employee\_Name** are also reflexive and thus trivial dependencies.

In essence, the reflexive rule is a basic logical inference: if you know the values of a set of attributes, you inherently know the values of any subset of those attributes. It's considered a **trivial functional dependency** if the right-hand side (dependent) is a subset of the left-hand side (determinant). Such dependencies are always satisfied by all instances of a relation.

## Q12 All NF explanation

Database normalization is the process of organizing data in a database to minimize redundancy and eliminate undesirable characteristics like insertion, update, and deletion anomalies. It involves dividing larger tables into smaller ones and linking them using relationships. The different normal forms represent stages of this process.

Here's an explanation of each normal form based on the sources:

- **First Normal Form (1NF):** A relation is in 1NF if **it contains only atomic values**. This means that an attribute (column) of a table cannot hold multiple values; it must hold only a single value. 1NF disallows multi-valued attributes, composite attributes, and their combinations. For example, a table with a column that stores multiple phone numbers for a single person is not in 1NF. To convert it to 1NF, each phone number should be in a separate row (potentially repeating other attribute values) or moved to a separate table with a foreign key linking back to the original entity.
- **Second Normal Form (2NF):** A relation is in 2NF if it meets two conditions:
  - It must be in 1NF.
  - **All non-key attributes are fully functional dependent on the primary key.** This means that every non-key attribute must depend on the entire primary key, not just a part of it. 2NF addresses partial dependencies. If a non-key attribute depends only on a proper subset of a composite primary key, the relation is not in 2NF. To achieve 2NF, such attributes and the part of the key they depend on are moved into a new table.
- **Third Normal Form (3NF):** A relation is in 3NF if it satisfies two conditions:
  - It must be in 2NF.
  - **It contains no transitive partial dependency.** A transitive dependency occurs when a non-key attribute depends on another non-key attribute. In other words, if  $X \rightarrow Y$  and  $Y \rightarrow Z$ , and both Y and Z are non-key attributes, then

there is a transitive dependency, and the relation is not in 3NF. To reach 3NF, the transitively dependent attributes are moved to a new table along with the attribute they depend on. A relation is in 3NF if for every non-trivial functional dependency  $X \rightarrow Y$ , either X is a superkey, or Y is a prime attribute (part of a candidate key). 3NF helps to further reduce data duplication and improve data integrity.

- **Boyce-Codd Normal Form (BCNF):** BCNF is an **advanced version of 3NF and is stricter**. A relation R is in BCNF if for every non-trivial functional dependency  $X \rightarrow Y$ , **X is a superkey**. If every determinant in a table is a superkey, then the table is in BCNF. BCNF addresses situations that 3NF might not, particularly when there are multiple candidate keys and dependencies between them. If a dependency exists where the determinant is not a superkey, the table needs to be decomposed to achieve BCNF. While BCNF aims for a higher level of normalization, it's not always possible to achieve BCNF while preserving all dependencies.
- **Fourth Normal Form (4NF):** A relation is in 4NF if it meets two criteria:
  - It must be in BCNF.
  - **It has no multi-valued dependencies.** A multi-valued dependency exists when for a single value of attribute A, there are multiple independent values of attribute B. This often occurs when an entity has multiple independent multi-valued attributes. To achieve 4NF, the table is decomposed into two or more tables, each containing one of the multi-valued dependencies.
- **Fifth Normal Form (5NF):** A relation is in 5NF if it satisfies two conditions:
  - It must be in 4NF.
  - **It does not contain any join dependency, and joining should be lossless.** 5NF, also known as Project-join normal form (PJ/NF), deals with situations where a table can be losslessly decomposed into three or more smaller tables based on join dependencies that are not implied by functional or multi-valued dependencies alone. 5NF aims to eliminate redundancy resulting from these complex dependencies by ensuring that the table cannot be further decomposed without loss of information. It's often satisfied when tables are broken down as much as possible to avoid redundancy.

The goal of normalization is to create database designs that are well-structured, efficient, and maintain data integrity by reducing redundancy and eliminating anomalies. Achieving higher normal forms generally leads to more tables and more joins to retrieve data but reduces data duplication and improves data consistency.

## Q13 CODDS rule

**Codd's Rules** are a set of **twelve rules** that **E.F. Codd** formulated after his extensive research on the Relational Model of database systems. According to him, a database must obey these rules to be considered a true relational database. These rules can be applied to any database system that manages stored data using only its relational capabilities. The first rule is considered a foundation for all the others.

Here's an explanation of each of Codd's twelve rules based on the provided excerpts:

- **Rule 1: Information Rule:** All data stored in a database, whether user data or metadata, **must be a value of some table cell**. Everything in a database must be stored in a table format.
- **Rule 2: Guaranteed Access Rule:** Every single data element (value) is **guaranteed to be logically accessible** by using a combination of **table-name, primary-key (row value), and attribute-name (column value)**. No other means, such as pointers, can be used to access data.
- **Rule 3: Systematic Treatment of NULL Values:** **NULL values** in a database must be given a **systematic and uniform treatment**. This rule is important because a NULL can be interpreted as one of the following: data is missing, data is not known, or data is not applicable.
- **Rule 4: Active Online Catalog:** The structure description of the entire database **must be stored in an online catalog**, known as the **data dictionary**, which can be accessed by authorized users. Users should be able to use the same query language to access the catalog that they use to access the database itself.
- **Rule 5: Comprehensive Data Sub-Language Rule:** A database can only be accessed using a **language having linear syntax** that supports **data definition, data manipulation, and transaction management operations**. This language can be used directly or through an application. If the database allows data access without this language, it's considered a violation. SQL (Structured Query Language) is mentioned as a computer language aimed to store and manipulate data in relational databases, with basic functions including inputting, modifying, and dropping data. Data Definition Language (DDL), Data Manipulation Language (DML), and Transaction Control Language (TCL) are parts of such a database language like SQL.
- **Rule 6: View Updating Rule:** All the **views** of a database that can theoretically be updated **must also be updatable by the system**.
- **Rule 7: High-Level Insert, Update, and Delete Rule:** A database must support **high-level insertion, updation, and deletion**. This should not be limited to single rows; it must also support operations like union, intersection, and minus to return sets of data records. DML (Data Manipulation Language) is used to retrieve, store, modify, delete, insert, and update data in the database, with examples like SELECT, UPDATE, and INSERT statements.

- **Rule 8: Physical Data Independence:** The data stored in a database **must be independent of the applications that access it**. Any change in the physical structure of a database must not impact how external applications access the data.
- **Rule 9: Logical Data Independence:** The logical data in a database **must be independent of its user's view (application)**. Any change in logical data (e.g., merging or splitting tables) should not affect the applications using it. This is noted as one of the most difficult rules to apply.
- **Rule 10: Integrity Independence:** A database must be **independent of the application** that uses it. All its **integrity constraints** can be independently modified without requiring changes in the application. This rule makes a database independent of the front-end application and its interface. Constraints, such as NOT NULL, UNIQUE, PRIMARY KEY, and FOREIGN KEY, are used by the Oracle server to prevent invalid data entry into tables and to prevent table deletion if dependencies exist. All constraints are stored in the data dictionary.
- **Rule 11: Distribution Independence:** The end-user **must not be able to see that the data is distributed** over various locations. Users should always have the impression that the data is located at one site only. This rule is considered the foundation of distributed database systems. Distributed database systems involve data residing at multiple sites.
- **Rule 12: Non-Subversion Rule:** If a system has an interface that provides access to low-level records, then the interface **must not be able to subvert the system** and bypass security and integrity constraints.

## Q14 Types of data structure, unstructured and semi structure

### Data Structures:

The sources mention **data structures** primarily in the context of how data is managed within a DBMS. For example, the **storage manager** in a DBMS implements **data structures** used to represent information stored on disk. These include:

- **Data files:** Where the database itself is stored.
- **Data dictionary:** Which stores metadata about the structure of the database.
- **Indices:** Which provide fast access to data items. Indexes store the value of specific fields, ordered by the value of the field. MongoDB, for instance, supports indexes on any field or sub-field of documents in a collection, including multikey indexes for arrays.

The **relational model** itself uses **tables** as a fundamental **data structure** to represent data and relationships. These tables have rows (records) and columns (attributes).

In the context of **NoSQL databases**, different **data structures** are employed depending on the type of NoSQL database. These include:

- **Key / Value stores**: Which use a simple **hash table** consisting of key-value pairs.
- **Column-oriented databases**: Where data is stored in sections of **columns**, logically grouped into column families.
- **Document databases**: Which use **document collections**, where individual documents (similar to JSON objects) can have multiple fields with dynamic schemas.
- **Graph databases**: Which use a **graph structure** with vertices (nodes) interconnected by edges to represent relationships.

### **Unstructured Data:**

**Unstructured data** is defined as data that **cannot be readily classified** and does not reside in a fixed field within a record or file. The sources provide several examples of unstructured data:

- Photos and graphic images.
- Videos.
- Streaming instrument data.
- Webpages.
- PDF files.
- PowerPoint presentations.
- Emails (body of the message and attachments).
- Blog entries.
- Wikis.
- Word processing documents.
- Books.
- Journals.
- Documents (like simple text documents).
- Metadata.
- Health records.
- Audio.
- Analog data.
- GPS Tracking information.

The sources note that a significant portion of the world's data, around **80%**, is unstructured. NoSQL databases with dynamic schemas are better suited for handling unstructured data compared to traditional SQL databases that require data to fit into tables.

### **Semi-Structured Data:**

**Semi-structured data** is described as a **cross between structured and unstructured data**. It has **some structure** but **lacks a strict data model structure**. Tags or other markers are used to identify certain elements within the data, but not all information has an identical structure. Examples provided in the sources include:

- Word processing software documents that include **metadata** like the author's name and creation date, while the main content is unstructured text.
- **Emails** that have fixed fields for sender, recipient, date, and time, along with the unstructured content of the message and any attachments.
- **Photos or other graphics** that can be tagged with keywords such as creator, date, location, and other terms, allowing for organization.
- **XML** and other **markup languages** are often used to manage semi-structured data.

The sources highlight that while semi-structured data has some organizational elements, not all entities of a particular type may have the same set of attributes. NoSQL databases, particularly document databases, can also effectively handle semi-structured data due to their flexible schemas.

## Q15 Explain 5v's of big data

The five V's of Big Data are **Velocity, Volume, Value, Variety, and Veracity**. According to the sources:

- **Velocity**: This refers to the **speed at which vast amounts of data are being generated, collected, and analyzed**. The rate of increase in emails, tweets, photos, and videos is very high. Big data technology allows for analyzing data as it is generated, enabling real-time access and analysis for applications like website performance monitoring and credit card verification.
- **Volume**: This refers to the **incredible amounts of data generated each second** from various sources such as social media, cell phones, and sensors. The sheer size of this data makes it difficult to store and analyze using traditional database technologies.
- **Value**: This refers to the **worth of the data being extracted**. Simply having large amounts of data is not useful unless it can be turned into valuable insights. There is a link between data and insights, but not all Big Data inherently holds value.
- **Variety**: This is defined as the **different types of data that can now be used**, including unstructured data. A significant portion of the world's data, around **80%**, is unstructured, encompassing photos, videos, and social media updates. Big data technologies enable the simultaneous handling of both structured and unstructured data.
- **Veracity**: This refers to the **quality or trustworthiness of the data**. It questions the accuracy and reliability of the data, considering issues like typos and abbreviations in social media posts.

## Q16 CAP theorem

The **CAP Theorem**, also known as Brewer's Theorem, states that it is **impossible for a distributed computer system to simultaneously provide all three of the following guarantees: Consistency, Availability, and Partition tolerance**. A distributed system can satisfy any two of these guarantees at the same time but not all three.

- **Consistency (C):** This means that **all nodes see the same data at the same time**. After an update operation, every client should see the same data, ensuring data consistency.
- **Availability (A):** This implies that **node failures do not prevent other survivors from continuing to operate**. It's a guarantee that **every request receives a response about whether it succeeded or failed**. The system should always be on, ensuring service guarantee availability. Basic Availability means the system will fulfill requests, even with partial consistency, and the database should appear to work.
- **Partition tolerance (P):** This means that **the system continues to operate despite arbitrary partitioning due to network failures** (e.g., message loss). The system continues to function even if the communication among the servers is unreliable.

According to the sources, different types of distributed database systems prioritize different pairs of these guarantees:

- **CA (Consistency and Availability):** RDBMS are mentioned in this category.
- **CP (Consistency and Partition tolerance):** MongoDB, HBase, and Redis are listed as examples.
- **AP (Availability and Partition tolerance):** Cassandra, CouchDB, DynamoDB, and Riak are given as examples.

The sources also touch upon related concepts in the context of prioritizing Availability over strong Consistency in some distributed systems:

- **Soft State:** Data storage may not contain the write-consistent state. Different replicas on different shards might not be mutually consistent at all times.
- **Eventual Consistency:** At some point in the future, data will converge to a consistent state. This is a delayed consistency, as opposed to the immediate consistency of the ACID properties. A BASE model normally focuses on availability and scaling but does not guarantee data consistency.

Therefore, when designing distributed applications, architects must make a trade-off and choose two out of the three CAP guarantees to focus on based on the specific requirements of their system

## Q17 mapReduce explanation with example

**MapReduce** is a **data processing paradigm for condensing large volumes of data into useful aggregated results**. In MongoDB, the mapReduce database command is used to perform these operations.

Here's a breakdown of how MapReduce works according to the sources:

- The mapReduce function **first queries a collection**.
- Then, it applies a **map function** to each resulting document. This JavaScript function **emits key-value pairs**. The emit(key, value) function associates a value with a key, and the map function can call this any number of times per input document. Inside the map function, the current document can be referenced using this.
- For keys that have multiple emitted values, MongoDB then applies a **reduce function**. This JavaScript function **collects and condenses the aggregated data** based on the keys. The reduce function takes the key and an array of values as arguments and should return a result. MongoDB will not call the reduce function for a key that has only a single value. Importantly, MongoDB can invoke the reduce function more than once for the same key.
- Finally, MongoDB **stores the results in a specified output collection**. The output location is specified using the out parameter.

An optional **finalize function**, which follows the reduce method and can modify the output. This JavaScript function receives the key and the reducedValue from the reduce function as arguments and should return a modifiedObject.

Furthermore, the mapReduce command has several optional parameters, including:

- query: Specifies selection criteria for documents.
- sort: Specifies sorting criteria.
- limit: Specifies the maximum number of documents to return.
- scope: Specifies global variables accessible in the map, reduce, and finalize functions.
- jsMode: Specifies whether to convert intermediate data into BSON format between the map and reduce functions (defaults to false).
- verbose: Specifies whether to include timing information in the result (defaults to true).

### Example:

Example of using mapReduce on a posts collection. The goal is to count the number of active posts for each user.



The map function emits the user\_name as the key and 1 as the value for each document where the status is "active":

```
function() { emit(this.user_name,1); }
```

The reduce function takes the key (user\_name) and an array of values (all 1s for posts by that user) and returns the sum of these values, effectively counting the posts:

```
function(key, values) {return Array.sum(values)}
```

The mapReduce command to achieve this is:

```
db.posts.mapReduce(
  function() { emit(this.user_name,1); },
  function(key, values) {return Array.sum(values)},
  {
    query:{status:"active"},
    out:"post_total"
  }
);
```

The output shows the name of the resulting collection (post\_total) and some statistics about the operation. To see the actual counts, you would then use the find() operator on the post\_total collection. This example demonstrates how mapReduce can process documents, group them by a key, and perform an aggregation (summing) to get a final result.

## Q18 log based recovery

Log-based recovery is a technique used in database management systems to ensure **database consistency, transaction atomicity, and durability despite system failures**. It involves maintaining a **log on stable storage** that records database update activities. The log is a sequence of **log records**.

When a transaction  $T_i$  starts, it writes a **<T\_i start> log record**. Before  $T_i$  executes a write(X) operation, a log record **<T\_i, X, V1, V2> is written**, where  $V1$  is the value of data item  $X$  before the write, and  $V2$  is the value to be written. This record indicates that  $T_i$  has performed a write on  $X$ , which had value  $V1$  and will have value  $V2$  after the write. When  $T_i$  finishes its last statement, a **<T\_i commit> log record** is written.

There are two main approaches to log-based recovery:

### 1. Deferred Database Modification:

- In this scheme, all log records are written to stable storage, but the **database is only updated after a transaction commits**.

- A write(X) operation results in a log record  $\langle T_i, X, V \rangle$  being written, where  $V$  is the new value for  $X$ . The old value is not needed in this scheme. The actual write to the database is deferred.
- When  $T_i$  partially commits,  $\langle T_i \text{ commit} \rangle$  is written to the log.
- During recovery after a crash, a transaction  $T_i$  needs to be **redone** if and only if both  $\langle T_i \text{ start} \rangle$  and  $\langle T_i \text{ commit} \rangle$  are present in the log. Redoing  $T_i$  involves setting the value of all data items updated by the transaction to their new values.

## 2. Immediate Database Modification:

- In this scheme, the **database is modified immediately after every write operation**, and each log record precedes the actual database modification.
- Because undoing might be necessary, **update logs must contain both the old and the new values**:  $\langle T_i, X, V_1, V_2 \rangle$ .
- The output of updated blocks can occur at any time before or after a transaction commits.
- The recovery procedure in this scheme involves two operations:
  - **undo( $T_i$ )**: Restores the value of all data items updated by  $T_i$  to their old values, working backward from the last log record for  $T_i$ .
  - **redo( $T_i$ )**: Sets the value of all data items updated by  $T_i$  to their new values, working forward from the first log record for  $T_i$ .
- Both undo and redo operations must be **idempotent**.
- During recovery after a failure:
  - A transaction  $T_i$  needs to be **undone** if the log contains  $\langle T_i \text{ start} \rangle$  but not  $\langle T_i \text{ commit} \rangle$ .
  - A transaction  $T_i$  needs to be **redone** if the log contains both  $\langle T_i \text{ start} \rangle$  and  $\langle T_i \text{ commit} \rangle$ .
- Undo operations are performed before redo operations during recovery.

To streamline the recovery process and avoid scanning the entire log, **checkpoints** are periodically performed. A checkpoint involves:

1. Outputting all log records currently in main memory to stable storage.
2. Outputting all modified buffer blocks to disk.
3. Writing a **<checkpoint> log record** to stable storage.

During recovery with checkpoints, the system only needs to consider the most recent transaction  $T_i$  that started before the checkpoint and any transactions that started after  $T_i$ . The recovery process involves scanning backward from the end of the log to find the most recent **<checkpoint>** record, then continuing backward to find a  $\langle T_i \text{ start} \rangle$  record. Only the part of the log following this start record needs to be considered. For immediate modification, transactions with  $\langle T_i \text{ start} \rangle$  but no  $\langle T_i \text{ commit} \rangle$  are undone, and transactions with both are redone. In some checkpoint schemes, the **<checkpoint L>** record includes a list  $L$  of active transactions at the time of the checkpoint. During recovery, undo and redo

lists are created based on the log records and the checkpoint information. Transactions in the undo list are undone, and those in the redo list are redone.

Log-based recovery ensures that even if the system crashes, the database can be restored to a consistent state by replaying or undoing the operations recorded in the log.

## Q19 acid property

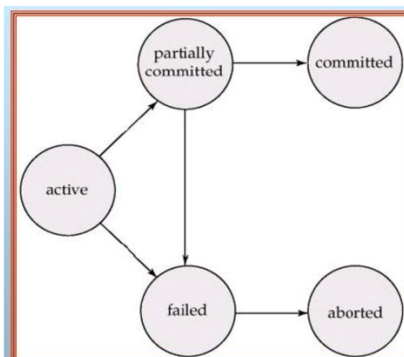
**ACID** is an acronym representing a set of properties of database transactions intended to guarantee validity even in the event of errors, power failures, etc.. The ACID properties are:

- **Atomicity:** This property ensures that **either all operations of the transaction are properly reflected in the database or none are**. The sources provide an example: when Steve transfers \$100 to Negan's account, this involves debiting Steve's account and crediting Negan's account. Atomicity means that if the system fails after debiting Steve's account but before crediting Negan's, the debit should be rolled back to maintain atomicity. In essence, a transaction is treated as a single, indivisible unit of work: "All or nothing". Transaction Control Language (TCL) commands like **COMMIT** (to persist changes) and rollback are used to manage this.
- **Consistency:** This property ensures that the **execution of a transaction in isolation preserves the consistency of the database**. The database starts in a consistent state, and a transaction moves it to another consistent state. For example, in a fund transfer, the consistency requirement might be that the sum of the balances of the two accounts involved remains unchanged after the transaction.
- **Isolation:** This property dictates that **although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions**. For every pair of transactions  $T_i$  and  $T_j$ , it should appear to  $T_i$  that either  $T_j$  finished before  $T_i$  started, or  $T_j$  started after  $T_i$  finished. This prevents one transaction from seeing the intermediate, uncommitted changes of another transaction, which could lead to inconsistencies.
- **Durability:** This property guarantees that **after a transaction completes successfully (commits), the changes it has made to the database persist, even if there are system failures**. Once a transaction is committed, the database system must ensure that the changes are permanent. Log-based recovery mechanisms are often used to achieve durability. Before a transaction  $T_i$  executes a write(X) operation, a log record  $\langle T_i, X, V1, V2 \rangle$  is written, where  $V1$  is the old value and  $V2$  is the new value. When  $T_i$  finishes, a  $\langle T_i \text{ commit} \rangle$  log record is written.

In contrast to SQL databases that emphasize ACID properties, NoSQL databases often trade off "ACID" for other properties like availability and partition tolerance, as described by Brewer's CAP theorem. Some NoSQL systems follow the BASE model (Basically Available, Soft-State, Eventually Consistent), which prioritizes availability and scaling but does not guarantee immediate data consistency. Eventual consistency means that data will converge to a consistent state at some point in the future, but there might be a delay.

## Q20 explain transaction state diagram

### Transaction State



The sources explain that a transaction progresses through different states during its execution. The **transaction state diagram** illustrates these states and the transitions between them. According to the sources, the typical states of a transaction are:

- **Active:** This is the **initial state** of a transaction while it is executing its operations. During this state, the transaction accesses and possibly updates data items.
- **Partially committed:** A transaction enters this state **after its final operation (last statement) has been executed**. At this point, all the operations have been performed, but the transaction has not yet been officially committed to the database. In deferred database modification, the actual database updates are performed after reaching this state.
- **Failed:** A transaction moves to the failed state **after the discovery that normal execution can no longer proceed**. This could be due to various reasons such as logical errors within the transaction or system errors like deadlocks.
- **Aborted:** After a transaction enters the failed state, it must be **rolled back**, meaning all the changes it has made are undone, and the database is restored to its state prior to the start of the transaction. Once the rollback is complete, the transaction is in the aborted state. The source mentions two options after a transaction is aborted: it can be restarted (only if there was no internal logical error) or it can be terminated (killed).
- **Committed:** A transaction reaches the committed state **after it has successfully completed all its operations and its changes have been permanently recorded**

**in the database.** The source notes that when a transaction finishes its last statement and is successful, a <Ti commit> log record is written. The TCL command **COMMIT** is used to persist these changes.

The transitions between these states are as follows:

- A transaction starts in the **active** state.
- From the **active** state, upon completion of the final operation, a transaction moves to the **partially committed** state.
- During the **active** or **partially committed** state, if a failure occurs, the transaction moves to the **failed** state.
- From the **failed** state, the system initiates a rollback, and upon successful undoing of the transaction's operations, the transaction enters the **aborted** state.
- If a transaction in the **partially committed** state successfully completes the commit protocol (making its changes durable), it transitions to the **committed** state.

The source highlights that for a transaction to see a consistent database, the database might be inconsistent during transaction execution, but it must be consistent once the transaction is committed. The **atomicity** property ensures that if a transaction fails before reaching the committed state, its operations are not reflected in the database (or are rolled back).

## Q21 consider following schedule conflict serializability or not and view serializability

We will be given a schedule which we have to determine whether it is conflict serializability or not and view serializability

## Q22 explain lock based and two phased protocol

Based on the sources, here's an explanation of lock-based and two-phase locking protocols.

### Lock-Based Protocols:

- A **lock** is a mechanism used in Database Management Systems (DBMS) to control **concurrent access to data items**. The goal is to manage simultaneous operations without conflicting with each other.
- Transactions must acquire a lock on a data item before accessing it and release the lock after accessing it.

- Data items can be locked in two modes:
  - **Exclusive (X) mode:** Allows a transaction to both read and write the data item. An X-lock is requested using the lock-X instruction. If a transaction holds an exclusive lock on an item, no other transaction can hold any lock on that item.
  - **Shared (S) mode:** Allows a transaction to only read the data item. An S-lock is requested using the lock-S instruction. Multiple transactions can hold shared locks on the same item simultaneously.
- Lock requests are made to the **concurrency-control manager**.
- A transaction is granted a lock if the requested lock is **compatible** with the locks already held on the item by other transactions. The lock-compatibility is typically defined by a matrix.
- If a lock cannot be granted, the requesting transaction is made to **wait** until all incompatible locks held by other transactions are released. The lock is then granted.
- However, simply using locking as described above is **not sufficient to guarantee serializability**. For instance, if a transaction reads two related data items, and another transaction modifies one of them in between the reads, the first transaction might get an inconsistent view.
- **Pitfalls of Lock-Based Protocols:**
  - **Deadlock:** A situation where two or more transactions are blocked indefinitely, waiting for each other to release locks. For example, transaction T3 might have a lock on item B and is waiting for a lock on item A held by T4, while T4 has a lock on item A and is waiting for a lock on item B held by T3. Deadlocks can be handled by rolling back one of the deadlocked transactions and releasing its locks. Deadlock can be described using a **wait-for graph**, where a cycle indicates a deadlock.
  - **Starvation:** A situation where a transaction is repeatedly denied a lock and cannot make progress, even though the data item it needs is being repeatedly locked and unlocked by other transactions. For example, a transaction requesting an exclusive lock might keep waiting while a series of other transactions acquire and release shared locks on the same item. Concurrency control managers can be designed to prevent starvation, for example, by prioritizing lock requests based on arrival time.

### Two-Phase Locking (2PL) Protocol:

- The Two-Phase Locking (2PL) protocol is a set of rules that all transactions must follow when requesting and releasing locks to **assure serializability**.
- The protocol has two phases:
  - **Growing Phase (Phase 1):** A transaction can acquire locks but **cannot release any locks**. The transaction acquires locks as needed.

- **Shrinking Phase (Phase 2):** A transaction can release locks but **cannot acquire any new locks**. Once a transaction releases a lock, it enters the shrinking phase and cannot request any more locks.
- A transaction initially starts in the growing phase. The transition to the shrinking phase occurs as soon as the first lock is released.
- **Benefits of 2PL:** This protocol **assures serializability**.
- **Drawbacks of 2PL:**
  - It **does not guarantee freedom from deadlocks**. Deadlocks can still occur if transactions acquire locks in a conflicting order during their growing phases.
  - **Cascading rollback** may occur under two-phase locking. If a transaction  $T_i$  reads a data item locked by another transaction  $T_j$ , and  $T_j$  later aborts, then  $T_i$  must also be rolled back (and any transactions that read data written by  $T_i$  would also need to be rolled back, and so on).

#### Variations of Two-Phase Locking to Avoid Cascading Rollback:

- **Strict Two-Phase Locking Protocol:** In addition to the two phases, this protocol requires that all **exclusive-mode locks** taken by a transaction be held **until that transaction commits or aborts**. This prevents other transactions from reading data written by an uncommitted transaction, thus avoiding cascading rollbacks.
- **Rigorous Two-Phase Locking Protocol:** This is even stricter than strict 2PL. It requires that **all locks (both shared and exclusive)** taken by a transaction be held **until the transaction commits or aborts**. With rigorous two-phase locking, transactions can be serialized in the order in which they commit. Most database systems implement either strict or rigorous two-phase locking.

In summary, lock-based protocols use locks to manage concurrent access to data, but they don't inherently guarantee serializability or freedom from issues like deadlock and cascading rollback. The two-phase locking protocol, with its growing and shrinking phases, ensures serializability but can still lead to deadlocks and cascading rollbacks. Strict and rigorous 2PL are variations that address the cascading rollback problem by holding exclusive (and all) locks until the transaction's completion.

## Q23 explain time stamp protocol

The sources explain that the **timestamp ordering protocol** is a concurrency control protocol that manages concurrent execution such that the **timestamps** assigned to transactions determine the **serializability order**.

Here's a breakdown of the key aspects of this protocol:

- **Timestamp Assignment:** Each transaction  $T_i$  in the system is associated with a unique, fixed timestamp, denoted by  $TS(T_i)$ . This timestamp is assigned by the database system before the transaction  $T_i$  starts its execution.
  - Timestamps are assigned in increasing order. If a transaction  $T_i$  has timestamp  $TS(T_i)$  and a new transaction  $T_j$  enters the system, then  $TS(T_i) < TS(T_j)$ .
  - Two common methods for implementing this are:
    - Using the value of the system clock when the transaction enters the system.
    - Using a logical counter that is incremented after each new timestamp assignment.
- **Data Item Timestamps:** For each data item  $Q$ , the protocol maintains two timestamp values:
  - **W-timestamp(Q):** This is the largest timestamp of any transaction that successfully executed a write( $Q$ ) operation.
  - **R-timestamp(Q):** This is the largest timestamp of any transaction that successfully executed a read( $Q$ ) operation.
- **Processing Read and Write Operations:** When a transaction  $T_i$  issues a read or write operation on a data item  $Q$ , the protocol checks the following conditions:
  - **read(Q) operation by  $T_i$ :**
    1. If  $TS(T_i) \leq W\text{-timestamp}(Q)$ , it means  $T_i$  needs to read a value of  $Q$  that has already been overwritten by a transaction with a later timestamp. Therefore, the read operation is **rejected**, and  $T_i$  is **rolled back**.
    2. If  $TS(T_i) > W\text{-timestamp}(Q)$ , the read operation is **executed**, and  $R\text{-timestamp}(Q)$  is updated to the maximum of its current value and  $TS(T_i)$ .
  - **write(Q) operation by  $T_i$ :**
    1. If  $TS(T_i) < R\text{-timestamp}(Q)$ , it implies that a transaction  $T_j$  with  $TS(T_j) > TS(T_i)$  has already read the value of  $Q$ . If  $T_i$ 's write were allowed, it would violate the timestamp order. Hence, the write operation is **rejected**, and  $T_i$  is **rolled back**.
    2. If  $TS(T_i) < W\text{-timestamp}(Q)$ , it means  $T_i$  is attempting to write an obsolete value of  $Q$ , as a transaction with a later timestamp has already written to  $Q$ . Thus, this write operation is also **rejected**, and  $T_i$  is **rolled back**.
    3. Otherwise (if  $TS(T_i) \geq R\text{-timestamp}(Q)$  and  $TS(T_i) \geq W\text{-timestamp}(Q)$ ), the write operation is **executed**, and  $W\text{-timestamp}(Q)$  is set to  $TS(T_i)$ .
- **Transaction Rollback and Restart:** If a transaction  $T_i$  is rolled back due to a rejected read or write operation, the system assigns it a **new timestamp** and **restarts** it.
- **Conflict Serializability:** The timestamp ordering protocol ensures **conflict serializability** because conflicting read and write operations are processed in timestamp order.



## Thomas' Write Rule: A Modification

The sources also mention **Thomas' Write Rule** as a modification to the basic timestamp ordering protocol. This rule aims to reduce unnecessary rollbacks in certain write operations.

Under Thomas' Write Rule, the rules for read operations remain the same. However, the rule for write operations is modified as follows:

1. If  $TS(T_i) < R\text{-timestamp}(Q)$ , the write operation is **rejected**, and  $T_i$  is **rolled back**.
2. If  $TS(T_i) < W\text{-timestamp}(Q)$ , this indicates an attempt to write an obsolete value. In this case, the write operation is **ignored**. This is the key difference from the basic protocol, where the transaction would be rolled back.
3. Otherwise (if  $TS(T_i) \geq R\text{-timestamp}(Q)$  and  $TS(T_i) \geq W\text{-timestamp}(Q)$ ), the write operation is **executed**, and  $W\text{-timestamp}(Q)$  is set to  $TS(T_i)$ .

Thomas' Write Rule avoids rolling back a transaction when it attempts to write a value that will never be read because a transaction with a later timestamp has already written to the same data item.

In summary, the timestamp ordering protocol uses timestamps to serialize transactions, ensuring that conflicting operations are executed in the order of their timestamps. Operations that would violate this order are rejected, leading to transaction rollback and restart. Thomas' Write Rule is an optimization that allows certain obsolete write operations to be ignored, potentially reducing the number of rollbacks.

## Q24 explain speedup and scaleup

### Speedup:

- **Speedup** refers to the concept of running a given task in **less time** by **increasing the degree of parallelism**.
- Imagine a task that takes a certain amount of time to execute on a smaller system. If we take the same task and run it on a system that is  $N$  times larger (meaning it has  $N$  times more resources, like processors), the goal is to reduce the execution time.
- The **speedup** achieved through parallelism is measured as the ratio of the execution time on the smaller machine ( $T_S$ ) to the execution time on the larger machine ( $T_L$ ), calculated as  $T_S/T_L$ .

- A parallel system is said to demonstrate **linear speedup** if the speedup achieved is equal to  $N$ , the factor by which the system size increased.
- If the speedup is less than  $N$ , the system is said to demonstrate **sublinear speedup**.

### Scaleup:

- **Scaleup** relates to the **ability to process larger tasks in the same amount of time by increasing the amount of resources**. It involves increasing both the size of the problem and the size of the system by a factor of  $N$ .
- Consider a task  $Q$  and a larger task  $Q_n$  that is  $N$  times bigger than  $Q$ . If the execution time of  $Q$  on a smaller machine is  $T_S$  and the execution time of  $Q_n$  on a parallel machine (that is  $N$  times larger) is  $T_L$ , then the scaleup is calculated as  $T_S/T_L$ .
- The parallel system is said to demonstrate **linear scaleup** on task  $Q$  if  $T_L$  is equal to  $T_S$ , meaning the  $N$  times larger task is completed in the same time as the original task on the smaller system.
- If  $T_L$  is greater than  $T_S$ , the system is said to demonstrate **sublinear scaleup**.

### Types of Scaleup:

- There are two main kinds of scaleup relevant in parallel database systems:
  - **Batch scaleup:** In this case, the **size of the database increases**, and the tasks are large jobs whose runtime depends on the size of the database. An example is scanning a relation whose size is proportional to the database size. The size of the database is the measure of the problem size.
  - **Transaction scaleup:** Here, the **rate at which transactions are submitted to the database increases**, and the size of the database increases proportionally to the transaction rate. This is relevant in transaction processing systems with small updates, like deposits or withdrawals, where transaction rates grow as more accounts are created. This type of processing is often well-suited for parallel execution as transactions can run concurrently and independently.

### Factors Limiting Speedup and Scaleup:

Speedup and scaleup are often sublinear due to several factors:

- **(A) Startup costs:** The overhead of starting up multiple processes for parallel execution can become significant, especially with a high degree of parallelism, and may overshadow the actual processing time.
- **(B) Interference:** Processes accessing shared resources like the system bus, disks, or locks can compete with each other, leading to waiting times and reduced efficiency.

- **(C) Skew:** It's often difficult to divide a task into perfectly equal-sized parts, leading to a skew in the workload distribution. The overall speed is then limited by the slowest part of the parallel execution. Increasing parallelism can sometimes increase the variance in the service times of parallel tasks.

## Q25 explain or elaborate parallel database architecture

A **parallel database system** consists of **multiple processors** and **multiple disks** connected by a **fast interconnection network**. These systems aim to improve performance through **parallelism**, which involves performing multiple operations simultaneously.

Two key performance measures for parallel systems are:

- **Throughput:** The number of tasks that can be completed in a given time interval.
- **Response time:** The time taken to complete a single task from submission.

Two important concepts related to the performance of parallel database systems are **speedup** and **scaleup**:

- **Speedup** refers to the ability to run a **fixed-sized problem in less time** by increasing the degree of parallelism. If a task takes time  $TS$  on a smaller system and time  $TL$  on an  $N$ -times larger system, the speedup is  $TS/TL$ . **Linear speedup** occurs when the speedup is equal to  $N$ . If the speedup is less than  $N$ , it's called **sublinear speedup**.
- **Scaleup** relates to the capacity to handle **larger tasks in the same amount of time** by increasing resources proportionally. If a task  $Q$  takes time  $TS$  and an  $N$ -times larger task  $Qn$  takes time  $TL$  on an  $N$ -times larger system, the scaleup is  $TS/TL$ . **Linear scaleup** is achieved if  $TL$  is equal to  $TS$ , meaning the larger task completes in the same time. If  $TL > TS$ , it's **sublinear scaleup**.

However, achieving linear speedup and scaleup can be limited by factors such as:

- **(A) Startup costs:** The overhead of initiating multiple processes can dominate computation time, especially with high parallelism.
- **(B) Interference:** Processes accessing shared resources (e.g., bus, disks) can lead to contention and reduced efficiency.
- **(C) Skew:** Uneven division of tasks can cause some processors to be idle while others are still working, limiting overall speed. Increased parallelism can also increase the variance in service times of parallel tasks.

The architecture of the interconnection network that connects processors and disks is also crucial and can include structures like a **bus**, **mesh**, or **hypercube**.

There are several main parallel database architectures:

- **Shared Memory:** In this architecture, all processors and disks have access to a **common main memory**. Communication between processors is very efficient as data in shared memory can be accessed by any processor. However, this architecture has limited **scalability** (typically up to 32 or 64 processors) because the bus or interconnection network becomes a **bottleneck** as the number of processors increases.
- **Shared Disk:** Here, each processor has its **own private memory**, but all processors can **directly access all disks** through an interconnection network. This overcomes the memory bus bottleneck of the shared-memory architecture and offers a degree of **fault tolerance**. If a processor fails, others can take over its tasks. However, the **interconnection to the disk subsystem can become a bottleneck**, especially with heavy disk access. Communication between processors is also slower compared to shared memory as it has to go through the network.
- **Shared Nothing:** In this model, each node in the system consists of its **own processor, memory, and one or more disks**. Processors at different nodes communicate via a high-speed interconnection network. This architecture is highly **scalable** and avoids the I/O bottleneck of the other architectures because local disk references are handled locally. The main drawbacks are the **higher costs of communication and nonlocal disk access**, as sending data involves software interaction at both ends. Examples of early systems using this architecture include Teradata, Grace, and Gamma.
- **Hierarchical:** This architecture is a **hybrid** that combines characteristics of shared-memory, shared-disk, and shared-nothing architectures. For instance, the top level might be a shared-nothing architecture where each node is a shared-memory system with a few processors. It allows for building systems with different levels of resource sharing to suit specific needs.

Commercial parallel database systems today utilize several of these architectures. The choice of architecture depends on factors such as the expected workload, scalability requirements, and cost considerations.

## Q26 explain data storage method -replication or fragmentation

There are two main approaches to storing a relation in a distributed database: **replication** and **fragmentation**. These methods can also be combined, where a relation is fragmented, and then each fragment is replicated.

### 1. Data Replication:

- **Definition:** Data replication involves maintaining **several identical replicas (copies)** of a relation and storing each replica at a different site. In the most extreme case, a copy is stored at every site in the system, which is called **full replication**.
- **Advantages:**
  - **Availability:** If a site containing a replica of a relation fails, the relation can still be accessed from another site. This allows the system to continue processing queries involving that relation despite the failure.
  - **Increased Parallelism:** If most accesses to a relation involve only reading, multiple sites can process queries involving that relation in parallel. Having more replicas increases the likelihood that the needed data will be found at the site where the transaction is executing, thus minimizing data transfer between sites.
- **Disadvantages:**
  - **Increased Overhead on Update:** The system must ensure that all replicas of a relation remain consistent. Therefore, whenever a relation is updated, the update must be propagated to all sites containing replicas, leading to increased overhead. For example, in a banking system, the balance of an account must be the same at all sites where the account information is replicated.

### 2. Data Fragmentation:

- **Definition:** Data fragmentation involves **partitioning a relation into several fragments** and storing each fragment at a different site. The fragments contain **sufficient information to reconstruct the original relation**. There are two main types of fragmentation:
  - **Horizontal Fragmentation:** Each **tuple (row)** of the relation is assigned to one or more fragments. A horizontal fragment can be defined as a **selection** on the global relation using a predicate. The original relation can be reconstructed by taking the **union** of all horizontal fragments. For instance, the account relation could be horizontally fragmented based on branch-name, with one fragment for each branch.
  - **Vertical Fragmentation:** The **schema (columns)** of the relation is split into several smaller schemas. All schemas must contain a **common candidate**

**key (or superkey)** to ensure the lossless join property, allowing the reconstruction of the original relation by taking the **natural join** of the vertical fragments. A special tuple-id attribute might be added to each schema to serve as a candidate key. For example, an employee-info relation could be vertically fragmented into employee-private-info (employee-id, salary) and employee-public-info (employee-id, name, designation).

- **Advantages of Fragmentation:**
  - **Horizontal Fragmentation:**
    - Allows **parallel processing** on different fragments of the relation.
    - Allows tuples to be located at the sites where they are **most frequently accessed**, minimizing data transfer.
  - **Vertical Fragmentation:**
    - Allows different parts of a tuple to be stored where they are **most frequently accessed**.
    - The tuple-id attribute enables **efficient joining** of vertical fragments.
    - Allows **parallel processing** on different parts of the relation.
- **Combination:** Horizontal and vertical fragmentation can be mixed, and fragments can be further fragmented or replicated to an arbitrary depth.

In summary, **replication** enhances data availability and read performance but increases the overhead of updates due to the need to maintain consistency across all copies.

**Fragmentation**, on the other hand, improves performance by allowing parallel processing and placing data closer to where it is frequently used, but it adds complexity to data retrieval as queries might need to access multiple fragments across different sites to get the complete relation. The choice between replication and fragmentation (or a combination of both) depends on the specific requirements of the distributed database system, such as the frequency of updates versus reads, the importance of data availability, and the need for performance optimization.

## Q27 diff homogeneous and heterogenous

Distributed databases can be classified as either homogeneous or heterogeneous. The key difference lies in the uniformity of the database management systems (DBMS) and schemas across the different sites.

Here's a breakdown of the differences:

- **Homogeneous Distributed Database:**
  - All sites in the system use **identical database management system software**.
  - The sites are **aware of one another**.
  - The sites **agree to cooperate** in processing users' requests.
  - They typically share a **common global schema**.
  - All sites run the **identical DBMS software**.
- **Heterogeneous Distributed Database:**
  - Different sites may use **different schemas**.
  - Different sites may use **different database management system software**.
  - The sites **may not be aware of one another**.
  - They may provide **only limited facilities for cooperation** in transaction processing.
  - The system may be composed of a **variety of data models**, for example, relational, object-oriented, and hierarchical models.

In essence, a **homogeneous distributed database** is a uniform system where all parts are the same and work together closely. Conversely, a **heterogeneous distributed database** is a more diverse system where different parts may operate independently with potentially limited coordination.

## Q28 2pc and 3pc

Both **Two-Phase Commit (2PC)** and **Three-Phase Commit (3PC)** are **commit protocols** used in distributed database systems to ensure the **atomicity of global transactions**. This means that a transaction that spans multiple sites must either commit at all sites or abort at all sites.

Here's an elaboration of each protocol:

### Two-Phase Commit (2PC):

- **Initiation:** A transaction  $T$  is initiated at a site  $S_i$ , and its coordinator is  $C_i$ . When  $T$  completes its execution at all participating sites, these sites inform  $C_i$ . Then,  $C_i$  starts the 2PC protocol.
- **Phase 1: Obtaining a Decision (Prepare Phase):**
  - The coordinator  $C_i$  writes a **<prepare T> record to its log** and **forces the log to stable storage**.
  - $C_i$  then sends a **<prepare T> message** to all sites where  $T$  executed (participants).
  - Upon receiving the **<prepare T> message**, each participating site determines if it is willing to commit its part of  $T$ .
  - If a site is **not willing to commit**, it writes a **<no T> record to its log** and sends an **<abort T> message** to  $C_i$ .
  - If a site is **willing to commit**, it writes a **<ready T> record to its log** and **forces the log (including all records related to T) to stable storage**. The transaction manager then replies with a **<ready T> message** to  $C_i$ . **Crucially, a site that sends a <ready T> message has promised to either commit or abort as instructed by the coordinator and must hold any locks acquired by the transaction.**
- **Phase 2: Recording the Decision (Commit/Abort Phase):**
  - When  $C_i$  receives responses from all participating sites, or if a timeout occurs after sending **<prepare T>**:
    - If  $C_i$  received **<ready T> from all participants**, it writes a **<commit T> record to its log** and **forces the log to stable storage**. It then sends a **<commit T> message** to all participants.
    - If  $C_i$  received **at least one <abort T> message** or a timeout occurred, it writes an **<abort T> record to its log** and **forces the log to stable storage**. It then sends an **<abort T> message** to all participants.
  - Upon receiving a **<commit T> message**, a participant writes a **<commit T> record to its log** and then commits the transaction locally. It may then send an **<acknowledge T> message** to the coordinator.
  - Upon receiving an **<abort T> message**, a participant writes an **<abort T> record to its log** and then aborts the transaction locally.



- **Handling Failures:**
  - **Failure of a participating site before sending <ready T>:** The coordinator assumes an <abort T> response.
  - **Failure of a participating site after sending <ready T>:** The coordinator proceeds with the commit protocol as normal, and the failed site will recover based on the log (it will have a <ready T> record and will wait for the coordinator's decision).
  - **Failure of the coordinator before making a decision:** Participating sites that have sent <ready T> become **blocked**, waiting for the coordinator to recover and issue a commit or abort. This is a significant drawback of 2PC.
- **Network Partition:** If the coordinator and participants are in different partitions, those not with the coordinator might incorrectly assume coordinator failure and take action. However, if they eventually rejoin, consistency should be maintained.

### Three-Phase Commit (3PC):

- 3PC is designed to **avoid the blocking problem of 2PC** under certain assumptions, including no network partitioning and at most  $K$  site failures. It introduces an extra phase to ensure that if a coordinator fails after participants have agreed to prepare, a new coordinator can still lead the transaction to a consistent commit or abort.
- **Phase 1: Obtaining Preliminary Decision (Prepare Phase):** This phase is identical to the first phase of 2PC, where the coordinator asks participants to prepare and they respond with <ready T> or <abort T>.
- **Phase 2: Recording the Preliminary Decision (Pre-Commit Phase):**
  - If the coordinator receives <ready T> from all participants, it writes a **<precommit T> record to its log** and forces it to stable storage.
  - The coordinator then sends a **<precommit T> message** to all participants.
  - Participants receiving <precommit T> write a **<precommit T> record to their logs** and send an **<acknowledge T> message** to the coordinator.
- **Phase 3: Recording the Decision (Commit/Abort Phase):**
  - Once the coordinator has received acknowledgements from all participants for the <precommit T> message, it writes a **<commit T> record to its log** and forces it to stable storage.
  - The coordinator then sends a **<commit T> message** to all participants.
  - Participants receiving <commit T> write a **<commit T> record to their logs** and commit the transaction locally.
- **Handling Failures (Coordinator Failure Protocol):** If the coordinator fails, a new coordinator ( $C_{new}$ ) is elected.  $C_{new}$  then determines the status of the transaction  $T$  by examining the logs of the participants. Based on the presence of <commit T>, <abort T>, <ready T>, or <precommit T> records,  $C_{new}$  can decide to commit or abort  $T$  or restart the 3PC protocol, aiming to avoid permanent blocking. For instance, if any participant is in the committed state, the transaction can be committed globally. If any is in the aborted state, it can be aborted. The

<precommit T> state allows for a commit even if the original coordinator failed after all participants were ready.

### Key Differences:

- **Number of Phases:** 2PC has two phases (prepare and commit/abort), while 3PC has three phases (prepare, pre-commit, and commit/abort).
- **Blocking:** 2PC can suffer from blocking where participants that have voted to prepare may have to wait indefinitely for the coordinator to recover and provide the final decision. 3PC aims to eliminate this blocking by introducing the pre-commit state, allowing participants to make progress even if the coordinator fails (under the assumption of no network partitions).
- **Assumptions:** 3PC relies on stronger assumptions than 2PC, particularly the absence of network partitioning. In the presence of network partitions, 3PC might still lead to inconsistencies.
- **Complexity:** The added phase in 3PC increases the complexity of the protocol compared to 2PC.

In summary, while 2PC is a widely used protocol for ensuring atomicity in distributed transactions, it has a blocking problem. 3PC attempts to address this limitation by adding a pre-commit phase, but it relies on more restrictive assumptions about the system.

## Q28 advantages over file system

Database Management System (DBMS) offers several significant advantages over traditional file-processing systems. These advantages address many of the limitations inherent in managing data through individual files.

Here are some key advantages of using a DBMS:

- **Reducing Data Redundancy:** File-based systems often contain multiple files where the same information may be duplicated in different locations. For example, in a college using file systems, a student enrolled in two courses might have their details stored twice, leading to wasted storage space. A DBMS, however, aims to **minimize this data redundancy** by providing a more centralized and structured way to store information.
- **Sharing of Data:** In a file-based system, data is often isolated within individual files, making it difficult for different applications or users to access and share related information. A DBMS facilitates **data sharing** among authorized users and applications.
- **Data Integrity:** Data redundancy in file systems can lead to **data inconsistency**. If a student's address is stored in multiple files and updated in only one, the data becomes inconsistent. A DBMS provides mechanisms to enforce **data integrity**, ensuring that the data remains accurate and consistent across the database.
- **Data Security:** Securing data from unauthorized access is challenging in file systems. For instance, preventing a student from accessing teacher payroll details is difficult to implement in file-based systems. A DBMS offers better **data security** features, allowing for the implementation of access controls and constraints.
- **Privacy:** Similar to security, managing data privacy is also enhanced with a DBMS, allowing for more controlled access to sensitive information.
- **Backup and Recovery:** Implementing reliable backup and recovery procedures is complex in file systems. A DBMS typically provides built-in features for **backup and recovery**, making it easier to restore data in case of system failures.
- **Data Consistency:** As mentioned earlier, data redundancy in file systems leads to inconsistency. By reducing redundancy, a DBMS helps to maintain **data consistency**.
- **Ease of Accessing Data:** Retrieving specific data from file systems often requires writing new application programs for each new requirement, which can be a tedious process. A DBMS provides **structured query languages (like SQL)** that allow users to access and manipulate data easily without the need for custom programs in many cases.
- **Data Isolation:** Because data in file systems is scattered across various files in different formats, integrating data from different sources can be difficult. A DBMS aims to reduce **data isolation** by providing a unified view of the data.

- **Atomicity:** Ensuring that a transaction is treated as a single, indivisible unit ("all or nothing") is difficult to achieve in file-processing systems. For example, if a system fails during a money transfer, a file system might leave the system in an inconsistent state where money is debited from one account but not credited to another. A DBMS provides mechanisms to ensure the **atomicity** of transactions.
- **Concurrency Control:** When multiple users access the same data concurrently in a file system, anomalies can occur where changes made by one user are lost due to changes made by another. A DBMS provides **concurrency control** mechanisms (like locking) to manage simultaneous access and prevent such anomalies.

In summary, a DBMS offers numerous advantages over file-processing systems by providing a structured and managed environment for data storage, retrieval, and manipulation, addressing issues related to redundancy, consistency, security, and concurrency.